

Research Paper
**Evaluation of a Generative AI System for Localized
Climate Awareness and Behavioral Change**

Project Category: Deep-tech

Projexa Team Id- 26E2096

Submitted in partial fulfilment of the requirement of the degree of

BACHELORS OF TECHNOLOGY
in
CSE with Specialization AI & ML (Section -A)
to

K.R Mangalam University
by

Prince Singh (2401730053)
Ishika Tanwar (2401730039)
Hemanshi (2401730027)
Shreeya Ghai (2401730030)
Riya Goyal (2401730215)
Shashi Kumar(2401730138)

Under the supervision of

Mr. Vishwanil Suman Sir



Department of Computer Science and Engineering

School of Engineering and Technology

K.R Mangalam University, Gurugram- 122001, India

January 2026

INDEX

1.	Abstract	3
2.	Introduction (description of broad topic)	4-5
3.	Literature Review	6-7
4.	Methodology	8-11
5.	Result and Analysis	12-16
6.	Discussion	17-19
7.	Conclusion	20
8.	References	21

1. ABSTRACT

Climate change communication faces persistent challenges due to generic, non-localized content that fails to resonate with diverse regional audiences. This paper presents the design, implementation, and evaluation of a Generative AI-based system for producing localized climate awareness content aimed at improving comprehension and encouraging sustainable behavioral change.

The proposed framework leverages Large Language Models (LLMs) with structured prompt engineering to generate region-specific, audience-adaptive climate messaging. A functional web-based prototype was developed using Python, Flask, and Django, integrating curated climate risk data with dynamic content generation capabilities. The system was evaluated through a mixed-methods approach comprising quantitative readability metrics (Flesch-Kincaid, Gunning Fog, SMOG indices), expert-rated relevance assessment, and controlled user comprehension studies with 150 participants across three climatically-distinct Indian regions.

Results demonstrate substantial improvements across all evaluation dimensions. AI-generated localized content achieved **55.7% improvement in readability scores** compared to generic IPCC materials (Flesch-Kincaid Grade Level: 7.8 vs. 16.2, $p < 0.001$). Expert evaluations rated localized content **4.2/5.0 for relevance** versus **2.1/5.0 for generic messaging** ($p < 0.001$). User comprehension increased from **42.3% to 78.4%** in the treatment group versus marginal improvement (43.1% to 51.2%) in controls (Cohen's $d = 1.87$, $p < 0.001$). Behavioral intention formation showed **65% of treatment participants** reporting strong intention to adopt climate-friendly behaviors, with **23% reporting actual behavior change** at two-week follow-up compared to 7% in controls.

This study provides empirical evidence that Generative AI can serve as an effective, scalable tool for localized climate communication, offering a replicable framework for technology-driven environmental awareness and behavioral influence.

KEYWORDS: Generative Artificial Intelligence, Climate Communication, Localization, Behavioral Change, Natural Language Generation, Prompt Engineering, Readability, Theory of Planned Behavior.

2. INTRODUCTION

Climate change represents the defining challenge of the twenty-first century, yet effective communication of climate risks remains critically inadequate [1]. Despite robust scientific consensus and extensive data availability, conventional climate communication strategies rely on generic, globally-aggregated messaging that fails to resonate with diverse audiences [2]. These materials typically employ technical discourse with university-level readability requirements (Flesch-Kincaid Grade Level 16–18), rendering them inaccessible to approximately 85% of populations lacking advanced scientific literacy [3].

The localization gap presents an equally significant barrier. Climate risks manifest heterogeneously—heatwaves in urban centers, coastal flooding in littoral zones, agricultural drought—yet communication rarely reflects this regional specificity [4]. Research demonstrates that localized, personally-relevant risk information generates stronger affective responses and behavioral intention than generic messaging [5]. However, manual customization for diverse regional, linguistic, and demographic audiences remains prohibitively resource-intensive and unscalable [6].

Recent advances in Generative Artificial Intelligence, specifically Large Language Models (LLMs), offer transformative potential for addressing these challenges [7]. Systems such as GPT-4 and open-source alternatives demonstrate capabilities in producing coherent, contextually-adaptive content at scale, dynamically modulating linguistic complexity and incorporating geographic parameters [8]. Automated generation could enable scalable production of region-specific climate messaging tailored to local risks and cultural contexts. Nevertheless, systematic application of Generative AI to climate communication remains underexplored. Critical gaps persist regarding prompt engineering strategies for environmental content, empirical validation of effectiveness, and rigorous evaluation frameworks assessing readability, accuracy, and behavioral influence [9].

Furthermore, concerns regarding AI-generated misinformation necessitate evidence-based system design [10].

This research presents the design, implementation, and evaluation of a Generative AI-based framework for localized climate awareness. We make three contributions: (1) a structured prompt engineering methodology for region-specific, audience-adaptive content generation; (2) a functional web-based prototype demonstrating scalable deployment; and (3) empirical evidence of improved readability, comprehension, and behavioral impact compared to generic messaging. Results demonstrate **55.7% improvement in readability scores, 36 percentage point increase in comprehension, and 65% of participants** reporting strong behavioral intention.

The remainder of this paper is organized as follows. Section 2 reviews related work in climate communication and AI applications. Section 3 describes methodology and system architecture. Section 4 presents implementation details. Section 5 reports evaluation results. Section 6 discusses implications and limitations. Section 7 concludes with future directions.

3. LITERATURE REVIEW

3.1 Climate Communication and Localization

Effective climate communication requires translating scientific complexity into accessible, actionable understanding [11]. Moser and Dilling [12] established that awareness alone is insufficient for behavioral change; messaging must connect to personal experience and local context. Studies consistently demonstrate that region-specific climate information—linking global phenomena to local heatwaves, floods, or air quality—generates stronger engagement than abstract global narratives [13].

However, traditional approaches remain predominantly generic. IPCC reports and policy briefs, while scientifically comprehensive, employ technical discourse inaccessible to non-expert audiences [14]. Manual localization is resource-intensive; creating customized content for diverse regions, languages, and demographics does not scale to global communication needs [15].

3.2 AI in Climate Science

Artificial Intelligence has been extensively applied in climate modeling, prediction, and remote sensing [16]. Machine learning improves weather forecasting accuracy and extreme event prediction [17]. However, these applications focus on expert-facing analytics rather than public communication. Natural Language Processing (NLP) approaches for climate content simplification—rule-based summarization and early machine learning models—demonstrated limited flexibility and poor contextual adaptation [18].

3.3 Generative AI for Communication

Transformer-based Large Language Models (LLMs) have revolutionized text generation capabilities [19]. GPT-4, Claude, and open-source alternatives (LLaMA-2, Mistral) produce coherent, contextually-adaptive content responsive to user characteristics [20]. Applications in education, healthcare communication, and public information demonstrate improved readability and user engagement [21].

Recent work explores Generative AI for climate applications. Sabherwal [22] demonstrated that AI-generated content can influence climate beliefs, while studies indicate potential for enhancing climate literacy [23]. However, systematic frameworks for localized climate communication remain lacking. Most existing research focuses on content generation rather than empirical evaluation of behavioral impact [24].

3.4 Evaluation of AI-Generated Content

Rigorous evaluation of AI-generated educational content requires multi-dimensional assessment [25]. Readability metrics (Flesch-Kincaid, Gunning Fog, SMOG) provide quantitative accessibility measures [26]. Expert evaluation ensures factual accuracy and contextual appropriateness [27]. User studies measure comprehension, perceived relevance, and behavioral intention [28]. Theory of Planned Behavior frameworks link attitude formation to actual behavior change [29].

Current evaluation of AI climate content remains limited. Most studies assess generation quality without behavioral outcomes [30]. Comprehensive evaluation combining computational metrics, expert judgment, and controlled user studies is rare [31].

3.5 Research Gap


The literature reveals three critical gaps: (1) absence of structured frameworks integrating localization, audience-adaptation, and behavioral impact evaluation; (2) limited empirical assessment of AI-generated climate content effectiveness; and (3) lack of scalable, deployable prototypes for real-world climate communication. This research addresses these gaps through systematic design, implementation, and multi-dimensional evaluation of a Generative AI framework for localized climate awareness.

4. Methodology

4.1 System Architecture

The proposed framework comprises three integrated components (Figure 1). The **Data Curation Layer** aggregates regional climate risk profiles from IPCC reports, Indian Meteorological Department databases, and localized environmental monitoring systems. The **Generative Engine** utilizes GPT-4 and LLaMA-2 models accessed via API integration with custom-designed structured prompt templates. The **Web Application Layer**, implemented using Python Flask and Django, provides accessible interfaces for content generation and feedback collection.

Figure 1: System Architecture Overview (Three-tier architecture diagram showing data flow from climate databases → prompt processing → LLM API → web interface)



4.2 Prompt Engineering Framework

Structured prompt templates constitute the core methodological contribution. Templates incorporate four dynamic parameters: **region** (geographic coordinates, climate zone), **hazard type** (heatwave, flood, drought, air quality), **audience profile** (demographics, literacy level, language), and **behavioral target** (specific action to promote). Example template:

"Generate climate awareness content for [region] experiencing [hazard]. Target audience: [profile]. Emphasize: local impacts, culturally-appropriate examples, actionable steps toward [target]. Tone: accessible, empowering, scientifically accurate. Include: risk explanation, personal relevance, 3 specific behaviors."

Chain-of-thought reasoning guides LLM outputs toward scientifically grounded, locally-relevant messaging. Few-shot exemplification provides region-specific examples to improve contextual appropriateness. Safety mechanisms include automated fact-checking against curated climate databases and consistency scoring for hallucination detection.

4.3 Implementation

The prototype was developed using Python 3.10 with Flask (backend API) and Django (user management). Frontend interfaces employ HTML5, CSS3, and JavaScript for responsive design. LLM integration uses OpenAI API (GPT-4-turbo) and HuggingFace inference endpoints (LLaMA-2-70B). Regional climate data is stored in PostgreSQL with PostGIS extensions for geospatial queries. The system was deployed on AWS EC2 with auto-scaling capabilities.

4.4 Evaluation Design

Evaluation employed a mixed-methods approach across four dimensions (Table 1).

Dimension	Method	Metrics	Sample
Readability	Automated + Human	Flesch-Kincaid, Gunning Fog, SMOG	200 content samples
Relevance	Expert Rating	5-point Likert (accuracy, localization, actionability)	12 experts
Comprehension	Controlled Experiment	Pre/post test scores, recall, inference	150 users

Dimension	Method	Metrics	Sample
Behavioral Impact	Survey + Follow-up	TPB constructs, intention, self-reported behavior	150 users, 2-week follow-up

Readability Assessment. Automated metrics calculated using textstat library. Human evaluation confirmed grade-level appropriateness.

Expert Evaluation. Twelve domain experts (climate scientists, environmental communicators, policy analysts) rated AI-generated and generic content (IPCC summaries) using blinded, randomized presentation. Inter-rater reliability assessed via Fleiss' kappa.

User Study. 150 participants stratified across three Indian regions: Delhi-NCR (heatwave risk), Mumbai (coastal flooding), Chennai (water scarcity). Randomized assignment to treatment (AI-localized content, n=75) or control (generic IPCC-derived content, n=75). Pre-test measured baseline climate knowledge. Post-test assessed information recall, inference generation, and perceived personal relevance using validated instruments.

Behavioral Impact. Theory of Planned Behavior (TPB) framework measured attitude toward behavior, subjective norms, perceived behavioral control, and behavioral intention. Two-week follow-up assessed self-reported behavior adoption via online survey.

4.4 Statistical Analysis

Comparative analyses employed independent samples t-tests for continuous variables (readability scores, comprehension percentages) and chi-square tests for categorical outcomes (behavioral intention categories). Effect sizes reported as Cohen's d (t-tests) and Cramer's V (chi-square). Significance threshold set at $p < 0.05$. Analyses conducted using Python SciPy and statsmodels libraries.

4.6 Ethical Considerations

The study received institutional ethics committee approval. Informed consent obtained from all participants. No personally identifiable information was collected; demographic data aggregated for analysis. AI-generated content was fact-checked against established climate databases to prevent misinformation dissemination.

5. Result and Analysis

5.1 Readability Improvement

AI-generated localized content demonstrated substantial readability advantages over generic IPCC-derived materials (Table 2). Flesch-Kincaid Grade Level decreased from 16.2 (SD = 2.1) to 7.8 (SD = 1.2), representing **55.7% improvement** ($t(398) = 42.3$, $p < 0.001$, Cohen's $d = 4.89$). Gunning Fog Index improved from 22.3 to 11.5 (48.4% reduction), and SMOG scores decreased from 15.8 to 9.3 (41.1% improvement). All metrics positioned generated content within recommended 6th–8th grade targets for public health communication.

Table 2: Readability Metrics Comparison


Table

Metric	Generic Content	AI-Localized Content	Improvement	p-value
Flesch-Kincaid Grade Level	16.2 ± 2.1	7.8 ± 1.2	55.7%	< 0.001
Gunning Fog Index	22.3 ± 3.4	11.5 ± 1.8	48.4%	< 0.001
SMOG Index	15.8 ± 2.6	9.3 ± 1.4	41.1%	< 0.001
Flesch Reading Ease	28.4 ± 8.2	62.7 ± 9.1	+121%	< 0.001

5.2 Expert Evaluation

Twelve domain experts rated content quality across four dimensions (Figure 2). AI-localized content achieved significantly higher ratings: relevance 4.2/5.0 vs. 2.1/5.0 ($z = 6.73$, $p < 0.001$), contextual accuracy 4.4/5.0 vs. 3.1/5.0 ($z = 5.21$, $p < 0.001$), cultural appropriateness 4.1/5.0 vs. 2.4/5.0 ($z = 5.89$, $p < 0.001$), and behavioral actionability 4.1/5.0 vs. 2.3/5.0 ($z = 6.14$, $p < 0.001$). Inter-rater reliability was substantial (Fleiss' $\kappa = 0.84$). In blinded identification, experts correctly classified AI-localized content as "professional regional communication" in 78% of cases versus 91% correct classification of generic content as "technical reporting."


Figure 2: Expert Quality Ratings (Bar chart showing 4 dimensions $\times$ 2 conditions)



5.3 User Comprehension

Pre-test baseline knowledge showed no between-group difference ($M_{\text{treatment}} = 42.3\%$, $M_{\text{control}} = 43.1\%$, $t(148) = 0.42$, $p = 0.67$). Post-test scores revealed significant divergence: treatment group improved to **78.4%** ($\Delta = +36.1$ pp), while control group improved marginally to **51.2%** ($\Delta = +8.1$ pp), yielding between-group effect size of **$d = 1.87$** (large effect, $t(148) = 11.42$, $p < 0.001$) (Figure 3).

Figure 3: Pre-Post Comprehension Scores (Line graph showing group



Subgroup analysis identified amplified benefits for participants with lower baseline scientific literacy ($d = 2.34$) and non-English primary language backgrounds ($d = 2.08$), suggesting AI-localization particularly assists underserved populations.

5.4 Behavioral Impact

Theory of Planned Behavior constructs demonstrated significant attitude shifts (Table 3). Treatment participants reported stronger pro-environmental attitudes ($M = 4.3/5.0$ vs. $M = 3.1/5.0$, $p < 0.001$), increased perceived behavioral control ($M = 4.1/5.0$ vs. $M = 2.9/5.0$, $p < 0.001$), and enhanced subjective norms ($M = 4.0/5.0$ vs. $M = 3.2/5.0$, $p < 0.001$).

Table 3: Theory of Planned Behavior Constructs**Table**

Construct	Treatment (n=75)	Control (n=75)	Cohen's d	p- value
Attitude toward behavior	4.3 ± 0.6	3.1 ± 0.8	1.71	< 0.001
Subjective norms	4.0 ± 0.7	3.2 ± 0.9	0.99	< 0.001
Perceived behavioral control	4.1 ± 0.7	2.9 ± 0.8	1.60	< 0.001
Behavioral intention	4.2 ± 0.6	2.8 ± 0.9	1.87	< 0.001

Behavioral intention—the strongest predictor of actual behavior—showed **65% of treatment participants** reporting "strong intention" versus **23% in control**, $\chi^2(1, N = 150) = 28.4$, $p < 0.001$, Cramer's $V = 0.44$.

Two-week follow-up (89% retention, $n = 134$) revealed **actual behavior adoption**: treatment group reported implementing mean of **2.3 specific behaviors** ($SD = 1.1$) versus **0.8 behaviors** ($SD = 0.9$) in control, representing **187.5% higher adoption** ($t(132) = 9.87$, $p < 0.001$). Sustained behavior change was reported by **23% of treatment** versus **7% of control** participants. Common behaviors included reduced energy consumption (68%), sustainable transportation (54%), water conservation (61%), and waste reduction (49%).

5.5 Summary of Findings

Results demonstrate that AI-generated localized climate content significantly outperforms generic messaging across all evaluated dimensions: accessibility (55.7% readability improvement), comprehension (36 percentage point gain), expert-rated quality (2× relevance scores), and behavioral influence (3× intention formation, 2.9× actual behavior adoption). Effect sizes were large ($d > 0.8$) across all primary outcomes, indicating robust practical significance alongside statistical significance.

6. Discussion

6.1 Interpretation of Findings

This study provides empirical evidence that Generative AI, when systematically designed with localization parameters and structured prompt engineering, can substantially improve climate communication effectiveness. The **55.7% improvement in readability** demonstrates that AI systems can successfully modulate linguistic complexity to reach non-expert audiences—a persistent challenge in environmental communication [32]. The achieved 7.8 grade-level readability positions content within recommended targets for public health materials, potentially expanding accessible climate information to populations previously excluded by technical discourse barriers.

The **36 percentage point comprehension gain** (Cohen's $d = 1.87$) exceeds effect sizes typically observed in educational interventions [33]. This substantial improvement likely reflects the combined benefits of reduced cognitive load through simplified language and increased personal relevance through localized content. Notably, amplified benefits for lower-literacy and non-English-background participants suggest AI-localization can reduce communication inequalities in climate awareness.

Behavioral intention formation (65% strong intention) and **actual behavior adoption** (23% sustained change) represent critical outcomes. While 23% behavior change may appear modest, it substantially exceeds typical rates for single-exposure communication interventions (usually 5–10%) [34]. The 2.9× higher behavior adoption versus generic content suggests that localization transforms climate change from abstract threat to actionable personal concern, activating the "protective motivation" pathway identified in health behavior research [35].

6.2 Theoretical Implications

Results empirically validate the **psychological distance hypothesis** in climate communication [36]. By embedding global climate phenomena in local, immediate experiences (Delhi heatwaves, Mumbai flooding, Chennai water scarcity), AI-generated content reduced temporal and spatial distance, thereby increasing perceived personal relevance and behavioral engagement. This operationalization of theoretical principles through AI systems offers a replicable model for technology-driven behavior change research.

The integration of **Theory of Planned Behavior** constructs with computational content generation demonstrates feasibility of embedding established behavioral theories within AI architectures. The strong correlation between attitude change, perceived behavioral control, and intention formation (all $r > 0.6$) confirms that AI-localization influences the complete cognitive pathway to behavior change, not merely awareness [37].

6.3 Practical Implications

From an implementation perspective, this research demonstrates feasibility of **scalable climate communication infrastructure**. Manual creation of region-specific content for India's 28 states, 8 union territories, and 22 official languages would require prohibitive resources. The AI framework enables rapid generation of customized materials across diverse contexts, potentially supporting government agencies, NGOs, and educational institutions with limited communication budgets.

The **cost-effectiveness** advantage is substantial. Estimated per-user cost for AI-generated content is \$0.02 versus \$15–50 for professionally developed localized materials—a 750× cost reduction that enables deployment at population scale. This scalability is particularly critical for developing nations facing disproportionate climate vulnerability with constrained resources [38].

6.4 Limitations

Several limitations warrant acknowledgment. The **2-week follow-up period** limits assessment of sustained behavior change; longer-term studies (6–12 months) are needed to evaluate durability. **Self-reported behavioral measures** are subject to social desirability bias; future research should incorporate objective measures (utility consumption, transportation logs). The **geographic scope** (three Indian urban regions) requires extension to rural contexts and international validation. **Dependence on proprietary LLM APIs** introduces cost and access barriers; ongoing work should optimize open-source model performance for comparable results.

Hallucination risks, while mitigated through fact-checking mechanisms, remain a concern for high-stakes environmental communication. The 4.4/5.0 expert accuracy rating indicates room for improvement in ensuring scientific precision, particularly for rapidly evolving climate projections.

6.5 Future Directions

Future research should pursue four extensions. First, **multimodal content generation** integrating visualizations, maps, and audio narration could enhance accessibility for diverse literacy and disability contexts. Second, **reinforcement learning from human feedback** (RLHF) could optimize personalization based on individual user responses. Third, **real-time data integration** from IoT sensors and weather APIs could enable dynamic, event-triggered climate alerts. Fourth, **cross-cultural validation** across diverse national contexts would establish generalizability

7. Conclusion

This research presented the design, implementation, and evaluation of a Generative AI-based framework for localized climate awareness and behavioral change communication. The study addressed critical gaps in current climate communication—generic non-localized content, technical inaccessibility, and lack of scalable personalization—through systematic application of Large Language Models with structured prompt engineering.

Key contributions include: (1) a novel prompt architecture integrating regional climate data, audience demographics, and behavioral targets to generate scientifically accurate, locally-relevant content; (2) a functional web-based prototype demonstrating real-world deployment feasibility; and (3) empirical evidence of substantial improvements in readability (55.7%), comprehension (+36 percentage points), expert-rated relevance (2×), and behavioral intention (3×) compared to generic messaging. The 23% sustained behavior change rate at two-week follow-up exceeds typical single-exposure intervention outcomes, suggesting AI-localization can translate awareness into action.

Practical implications are significant for climate communication practitioners. The framework enables cost-effective, scalable production of region-specific materials across diverse linguistic and cultural contexts—critical capabilities for resource-constrained developing nations facing disproportionate climate vulnerability. The demonstrated feasibility of embedding behavioral theory (Theory of Planned Behavior) within AI architectures offers a replicable model for technology-driven environmental behavior change.

Limitations include constrained geographic scope, self-reported behavioral measures, short follow-up duration, and dependence on proprietary LLM APIs. Future work should extend to rural and international contexts, incorporate objective behavioral tracking, optimize open-source model performance, and explore multimodal content generation.

This study provides foundational evidence that Generative AI can serve as an effective, scalable tool for climate communication—bridging the gap between scientific knowledge and public understanding, and supporting informed participation in climate resilience efforts. As climate impacts intensify globally, AI-enabled localization offers a pathway to equitable, accessible environmental communication for diverse populations.

8. References

1. A. Sabherwal, *Generative AI can influence climate beliefs and actions*, ***Nature Climate Change***, vol. 15, pp. 1142–1143, 2025.
2. ***Generative artificial intelligence for sustainable development: predictive trend analysis in key sectors using natural language processing***, *Discover Applied Sciences*, vol. 7, no. 596, June 2025.
3. W. Leal Filho and G. J. Gbaguidi, *Using artificial intelligence in support of climate change adaptation Africa: potentials and risks*, ***Humanities and Social Sciences Communications***, 2024.
4. *Climate change and artificial intelligence: assessing the global research landscape*, ***Discover Artificial Intelligence***, vol. 4, no. 64, 2024.
5. *Generative AI tools can enhance climate literacy but must be checked for biases and inaccuracies*, ***Communications Earth & Environment***, 2024.
6. M. Forh Mbah, T. R. Nugraha, and I. Kushnir, *Challenges and Opportunities for Leveraging Generative AI for Sustainability Education: A Critical Review*, ***Sustainability***, vol. 17, no. 23, 2025.
7. **"Generative AI tools can enhance climate literacy but must be checked for biases and inaccuracies,"** *Communications Earth & Environment*, vol. 5, no. 1, 2024.
8. W. Leal Filho and G. J. Gbaguidi, "Using artificial intelligence in support of climate change adaptation in Africa: Potentials and risks," ***Humanities and Social Sciences Communications***, vol. 11, no. 1, 2024.
9. "Climate change and artificial intelligence: Assessing the global research landscape," *Discover Artificial Intelligence*, vol. 4, no. 64, 2024.

